



GitLab

Pull Mirroring knowledge sharing session

Create Deep Dive
Tiago Botelho - Backend Engineer
18th of December 2018



- Share my knowledge of the pull mirroring feature with the entire GitLab team
- Make this “deep dive session” a reference for everyone that might need to work with pull mirroring in the future



- Demo
- Workflow
 - High level lifecycle
 - Architecture
- Code dive
- Troubleshooting
 - Debugging Redis
 - Debugging Sidekiq workers
 - Useful methods
- Useful Links
- Questions

We will not talk about



- Mirroring through SSH
- Push mirroring
- Bi-directional mirroring

What is pull mirroring?



- Feature available in GitLab Starter/Bronze tier
- Automatically pulls changes from an external repository into a project in GitLab
- Makes an effort to keep healthy mirrors synchronized with the external repository every 30 minutes
- A user is also able to update more often by using the “Update now” functionality
 - Also handles common failure scenarios gracefully
- Very useful for teams that have a canonical version of their code in an external repository and want to have a secondary version hosted on either GitLab.com or their own GitLab instance
 - E.g: Users have their code hosted on an external code hosting service
 - They want to leverage our CI service
 - They set up a pull mirror that is kept in sync and runs all the pipelines that were configured for that project



—





- There are a few core concepts that we need to explain in order for pull mirroring to make sense
 - Capacity
 - State transitions
 - State Management
 - Determining when a mirror update should be attempted again
- Key Metrics
 - Over 50k mirrors on GitLab.com
 - All of which were updated within the last 30 minutes



- Redis Set
- Contains the IDs of projects that are about to or currently are being updated by a Sidekiq worker
- The total capacity is a fixed number that can be configured by the GitLab instance admin
- It is used as a way of limiting the amount of mirrors that get added in the Sidekiq queue
- The objective is to always fill that capacity with as many mirrors as we can
 - This way the workers will always have work to perform
 - Translates into more frequent updates
- It should be a number higher than the configured Sidekiq concurrency
 - Making the capacity a lot higher than the concurrency that Sidekiq enables won't make a difference and will just translate into a bigger Sidekiq queue
- Making the value lower than the Sidekiq's concurrency will just translate into less frequent updates



- A mirror can be in one of the following five states:
 - None
 - Scheduled
 - Will be responsible for scheduling a worker to update the mirror
 - Started
 - Flags the time that the mirror started the update
 - Finished
 - Marks the time the mirror successfully finished
 - Will set the time when the mirror will get updated again
 - Failed
 - Marks the time the mirror finished unsuccessfully
 - Will increase the retry counter and set the time to update again
- The state machine is also useful to look for mirrors that are in inconsistent states
 - E.g: Mirrors in started state that don't have a running Sidekiq job



- There are three focal points that track the progress of the mirroring for each project
 - The Database
 - Holds information such as current status, job id, etc
 - Sidekiq
 - Provides the information about the mirroring queue
 - We are also able to know the status of each job in specific
 - Redis capacity set
 - Has the project IDs that are either in Sidekiq's queue or already being performed by one of the workers
- Spreading the information about which projects are currently getting updated helps the service become self-healing in some scenarios
 - Example: If a project in the Database says it has started
 - We can check if that job ID (stored in the Database) is still being performed or if it has finished already
 - This will tell us if Sidekiq was able to gracefully communicate with the DB in order to transition the project onto it's next stage

When should a mirror get scheduled?



- In order to determine when a mirror should be updated a formula was developed:

```
base_delay = (BACKOFF_PERIOD + rand(JITTER)) * (now() - last_update_started_at)
```

```
def set_next_execution_timestamp
  timestamp = Time.now
  retry_factor = [1, self.retry_count].max
  delay = [base_delay(timestamp), ::Gitlab::Mirror.min_delay].max
  delay = [delay * retry_factor, ::Gitlab::Mirror.max_delay].min

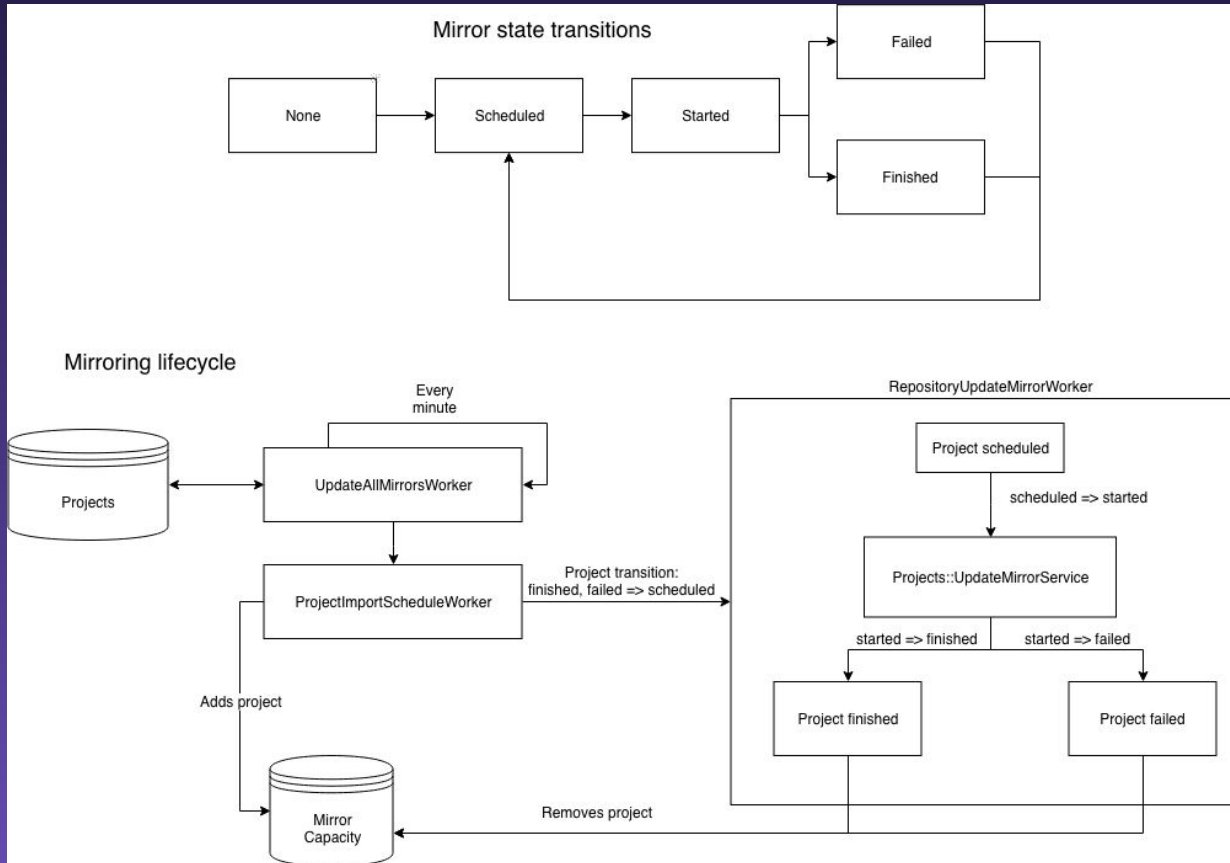
  self.next_execution_timestamp = timestamp + delay
end
```

- We want to penalise mirrors that fail often from running as frequently as healthy mirrors
- If a mirror reaches the maximum amount of retries, it will transition into a hard failed state where it won't get scheduled until a user takes action and solves the issue



1. The scheduler worker will pick all the mirrors that have the next execution time $< \text{now}()$
2. It will schedule mirrors until there is no more capacity available or when there are no more mirrors ready to be updated at the moment
3. After mirroring starts:
 - a. Fetch the changes from the provided remote URL
 - b. Update the respective branches with the new information
4. After updating the mirror:
 - a. Remove the project from the capacity list
 - b. Set the next execution time
 - c. Mirror finishes
 - i. Clear retry counter
 - d. Mirror fails
 - i. Retry counter gets incremented

Architecture





Questions?



Code Dive

Seeing what is behind the curtains



Questions?



- Example of an unhealthy mirroring system ([link](#))





- Always refer to the “project_mirror_data” table or the “ProjectImportState” model to check the state of your mirrors such as:
 - last_error
 - retry_count
 - jid
 - last_updated_at
 - last_successful_update_at
 - next_execution_timestamp
- The DB table is called “project_mirror_data” for legacy reasons even though ProjectImportState is used jointly by imports and forks as well



- Checking the available capacity
 - `Gitlab::Mirror.available_capacity`
 - Helps us debug situations where we might not be removing projects from the capacity
- `Project.mirrors_to_sync(Time.now)` will return all the mirrors ready to be picked for an update
 - Along with `Gitlab::Mirror.available_capacity` we are able to see if we have enough mirrors to completely fill the capacity up



- Check the status of the workers for each mirror in the scheduled/started state
 - `ProjectImportState.with_status([:scheduled, :started]).where.not(jid: nil).select(:jid)`
 - `Gitlab::SidekiqStatus.job_status(jids)`
- Retrieve the project IDs that are currently in the Redis set
 - `Gitlab::Redis::SharedState.with { |r| r.smembers(Gitlab::Mirror::PULL_CAPACITY_KEY) }`
 - Useful to look for projects that are stuck or with inconsistent information
 - Example: A finished/failed project ID should never be in that list



- Clear data inconsistencies
 - When the Database is inconsistent with Sidekiq
 - A project is started in the DB, but Sidekiq already considers it finished
 - `StuckImportJobsWorker` will look at the job ids maintained by the `Gitlab::SidekiqStatus` Redis key
 - Usually a timeout is the main cause for this scenario
 - When the capacity set is inconsistent with the DB and Sidekiq
 - A project is finished but the project id is still present in the capacity set
 - The only solution might be to remove that project ID from the capacity set
 - This is currently done manually (`StuckImportJobsWorker` will handle this in the future)
 - `Gitlab::Redis::SharedState.with { |redis| redis.del(Gitlab::Mirror::PULL_CAPACITY_KEY) }`
 - Only use this when the capacity is completely blocked!
 - `Gitlab::Mirror.decrement_capacity(project_id)`
 - This is currently done manually (`StuckImportJobsWorker` will handle this in the future)



- A healthy mirroring system





- [Pulling from a remote repository documentation](#)
- [StateMachine ActiveRecord module documentation](#)
- [Infrastructure Pull Mirroring Troubleshooting Guides](#)
 -



Questions?



Thank you

Tiago Botelho - Backend Engineer
tiago@gitlab.com